



LINUX  
**SECURITY**  
SUMMIT  
NORTH AMERICA

# Binding TDISP & Platform Attestation Reports for Confidential VMs

Anna Trikalinou, Microsoft

# Agenda

- Confidential Computing 101
- TDISP
  - I/O in CC without TDISP
  - I/O in CC with TDISP
- Attestation
  - Platform Attestation Evidence
  - TDISP Device Attestation Evidence
- Binding Evidence
  - Proposal
  - Open questions

# Confidential Computing 101

## EXISTING ENCRYPTION

### Data at rest

Encrypt inactive data when stored in blob storage, database, etc.

### Data in transit

Encrypt data that is flowing between untrusted public or private networks

## CONFIDENTIAL COMPUTING

### Data in use

Protect/encrypt data that is in use, while in RAM, and during computation



#### Insider threat

privileged admins abusing rights



#### Protect against

#### Hackers

exploiting bugs in the Hypervisor/OS



#### Third parties

accessing data without customer consent



The **protection of data in use** by performing computation in a hardware-based, attested Trusted Execution Environment (TEE).

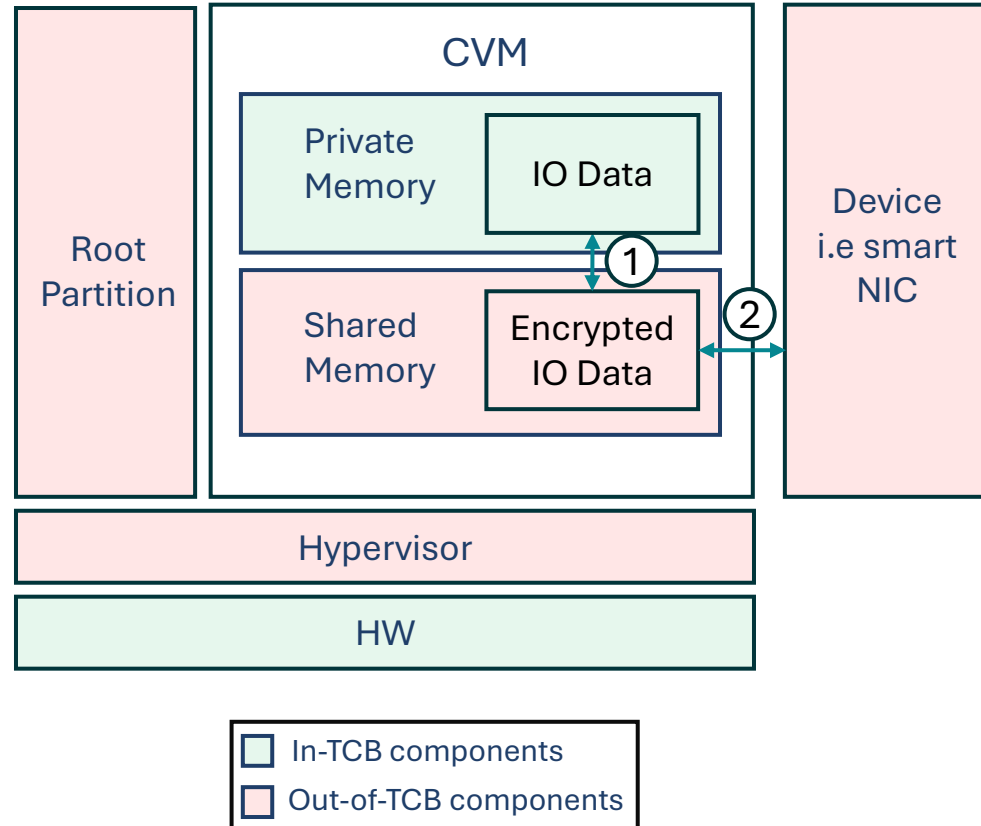
Verifiable assurance for data integrity, code integrity, and data confidentiality

# TEE Device Interface Security Protocol (TDISP)

- TDISP is a specification written from PCI SIG
- Defines how to:
  - Establish trust between CVM and a device (device attestation)
  - Secure the interconnect between the host and the device
  - Attach/detach a device interface to/from a CVM
- To support TDISP, we need:
  - Device changes, HW & FW
  - Host CPU changes (1st generation w. TDISP support: Intel GNR, AMD Turin)
  - Hypervisor, Host OS, CVM changes
  - Lots of infrastructure changes

# I/O in Confidential VM without TDISP (1/2)

- In CVM arch., memory is divided into:
  - **Private memory**, which is only accessible to in-TCB components (i.e. CVM and trusted HW)
  - **Shared memory**, which can be accessed by out-of-TCB components (i.e. HV, Host, devices)
- Without TDISP, the device cannot be added to the CVM TCB and can only access the CVM's shared memory.
- To do I/O, i.e. smart NIC:
  1. The device driver encrypts(/decrypts) & copies I/O data from private memory to shared (/from shared memory to private)
  2. Device accesses encrypted data via shared memory



# I/O in Confidential VM without TDISP (2/2)

- Other types of I/O devices:
  - GPU:
    - GPU needs to perform computations on the plaintext confidential data (i.e. training); cannot process encrypted data
    - CC principle is that only in-TCB components can access plaintext confidential data
    - NVIDIA Hopper offers a Confidential Computing architecture without TDISP to add the GPU into the Platform TCB and allow the GPU to work on confidential data
  - Crypto Accelerator:
    - Needs to process plaintext data; cannot double-encrypt
    - Will also require a custom Confidential Computing architecture to add the accelerator into the Platform TCB

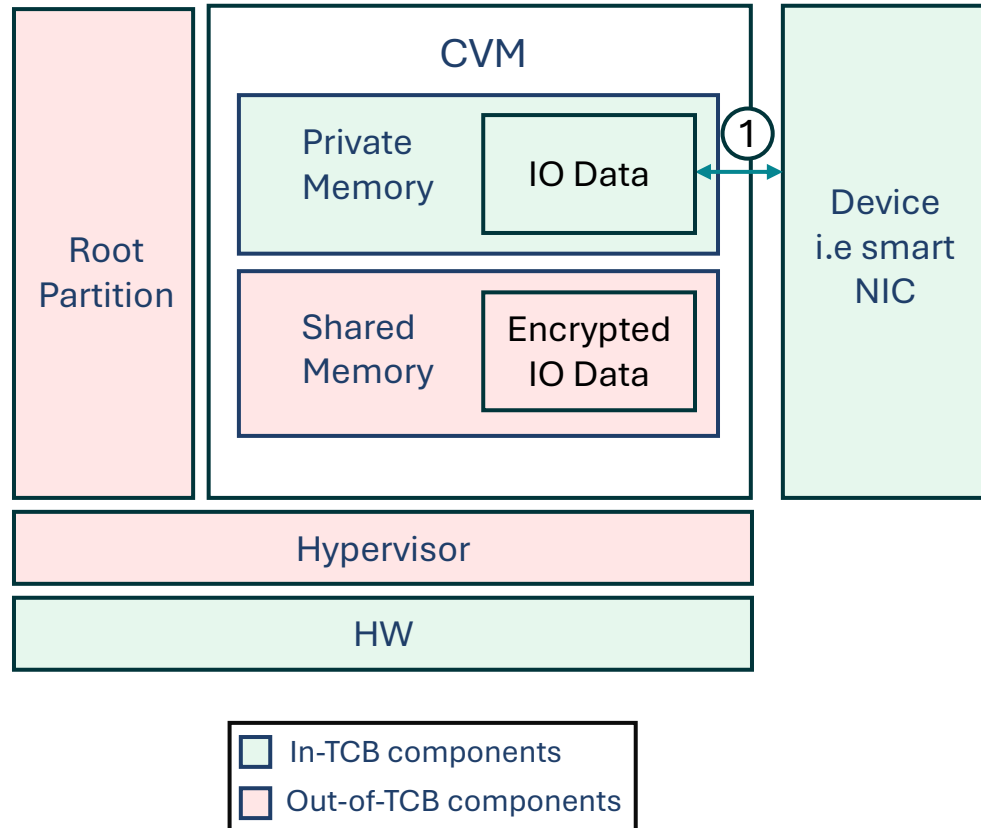
# I/O in Confidential VM without TDISP - Impact

- For smart NIC type of devices:
  - CPU encrypts/decrypts all data exiting/entering the TEE
    - Consumes CPU cycles, instead of working on CVM workload
  - Can impact IOPS and reduce performance
  - Can increase cost
- For GPU & Crypto Accelerator type of devices:
  - Need to define a TDISP-like Confidential Computing architecture to add the device into the CVM's TCB, so that the device can process confidential data

**Impact: Confidential Computing becomes more complex, less performant and more expensive than General Purpose Computing**

# I/O in Confidential VM with TDISP

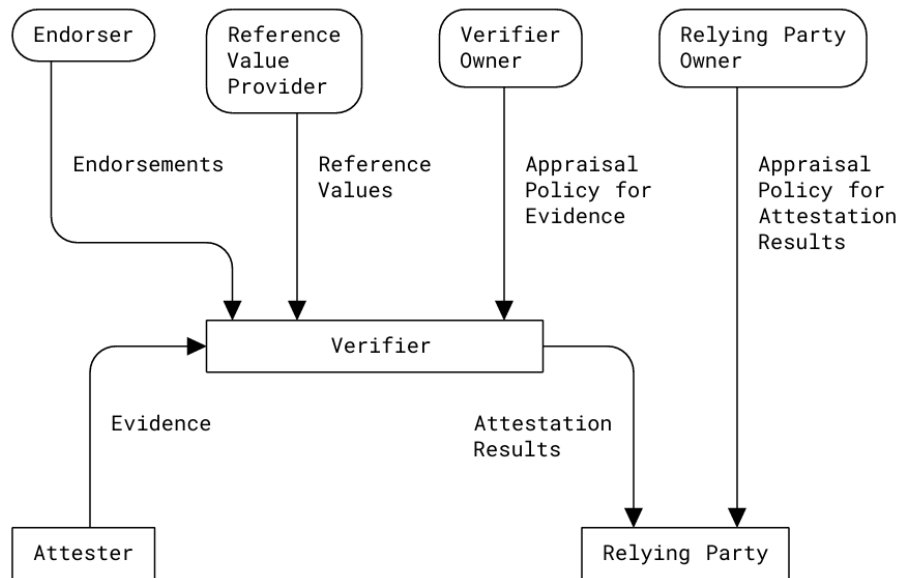
- With TDISP, the CVM can do device attestation and choose to add the device into the CVM's TCB
- An accepted device can access the CVM's private memory
- To do I/O, i.e. smart NIC:
  1. The accepted device can access the I/O data directly from the CVM's private memory. No CPU encryption/decryption is needed.
- **CVM performance is closer to GP**
  - Impact of IDE and other non-TDISP CVM features, i.e. memory encryption & integrity, IPI.
- **Less cost and complexity to support other accelerators**



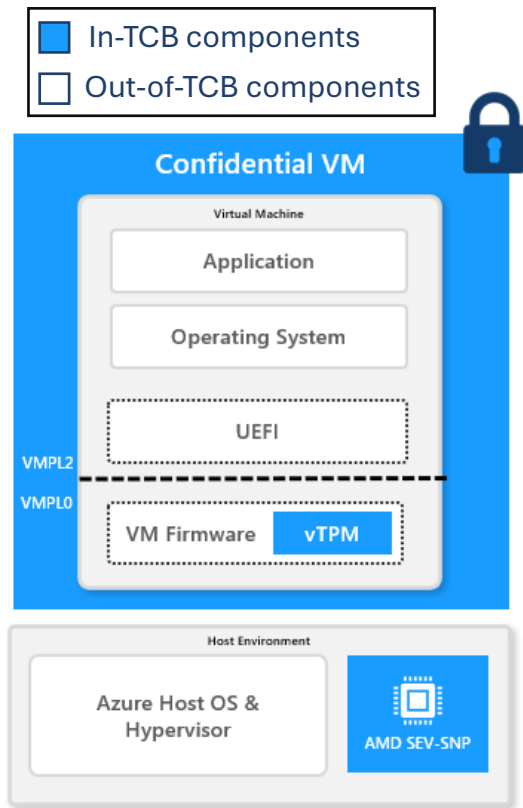


# Attestation in Confidential Computing

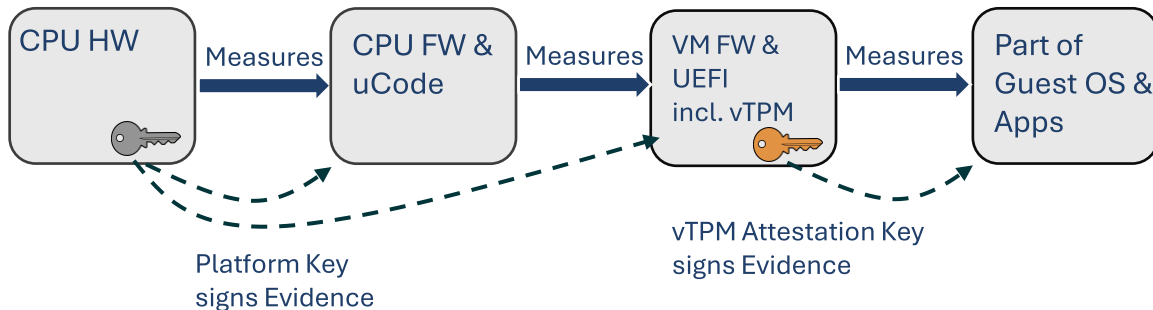
- CC offer verifiable assurance via Attestation
  - “Attestation is the process by which an **Attester** produces **Evidence** (believable, verifiable information about its own state), which is then evaluated by a **Verifier**. The Verifier appraises this Evidence using trusted inputs like **Endorsements** and **Reference Values**, and produces an **Attestation Result**. This result is then used by the **Relying Party** to make a trust decision about the Attester.” according to Remote ATtestation procedureS (RATS) Architecture
  - Platform Attestation and TDISP Attestation are distinct



# Platform Attestation Evidence

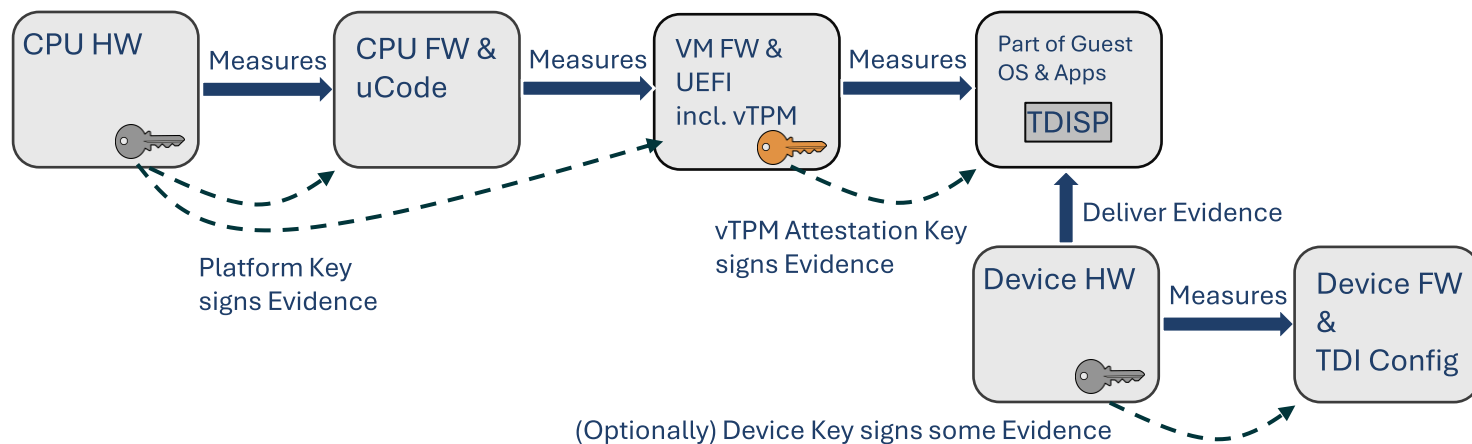


- Platform Attestation Evidence describe the in-TCB components of the CVM
  - Uses a chain of trust where one component measures the next and the root-of-trust is the CPU HW
  - A Verifier can check that the Guest Evidence is bound to the CPU Platform key, which proves that the CVM is running in a Confidential environment.
  - Guest kernel & drivers are measured at boot. Runtime attestation relies on IMA & systemd and their policy.
  - Contains no TDISP device information



# TDISP Attestation Evidence (1/2)

- TDISP Device Attestation Evidence describe the in-TCB components of a particular TDISP device
- The TDISP Device Attestation Evidence contain:
  - Device certificates, Measurements & Device Interface Report
- The Evidence is transferred from the Device to the CVM via the CPU FW/HW.
  - Since the CVM (VM FW or Guest OS) collects the Platform and TDISP device evidence, the CVM can verify that both evidence are mapped to the CVM itself.



# TDISP Attestation Evidence (2/2)

- **However, a remote Verifier cannot know if the TDISP Device Evidence originated from the same CVM as the Platform Evidence**
  - No existing binding between Platform and TDISP Device Evidence
  - An attacker could replay old TDISP Attestation Evidence with a CVM that doesn't have that TDISP device and cause a Relying Party to disclose secrets to the CVM that wouldn't otherwise
  - i.e. A Relying Party may have a policy to release proprietary training data to a CVM that has an attached TDISP GPU device.

# Binding Evidence – Goal

- Goal: Binding method is the same across OSES and CSPs
  - Becomes easier for anyone to implement their own Verifier ☺
- Proposed Solution:
  - CVM collects TDISP Device Attestation Evidence and uses evidence to extend a vTPM NVIndex
  - NVIndex gets signed using vTPM Attestation Key, which is already part of Platform Attestation Report

# TPM NVIndex 101

- TPM 2.0 Spec. defines that each TPM/vTPM implementation needs to have at least 6 KB of Non-Volatile Storage
  - Stores TPM-internal data (i.e. hierarchy seeds, counters, clocks, etc.)
  - Stores user-defined data and accessed via an NVIndex
- NVIndex is a handle (aka pointer) to a specific NV memory region. It includes:
  - Size of data
  - Attributes
    - Define static behavior and access rules of the NVIndex.
    - Defined at NVIndex creation time and cannot be changed afterwards.
    - i.e. If NV data acts as a counter, raw data or PCR, which hierarchy can read/write, etc.
  - Access control policy
    - Define dynamic authorization rules on who can access the NVIndex
    - i.e. require a password, key, specific PCR value to access NVIndex

# Binding Proposal (1/3)

- Define a PCR-like NVIndex that gets extended with the hash of each TDISP Device Attestation Evidence
- Create a file to store Event Log Records related to this NVIndex
- Every time a new TDISP Device Attestation Evidence is retrieved by the CVM, the CVM would do:

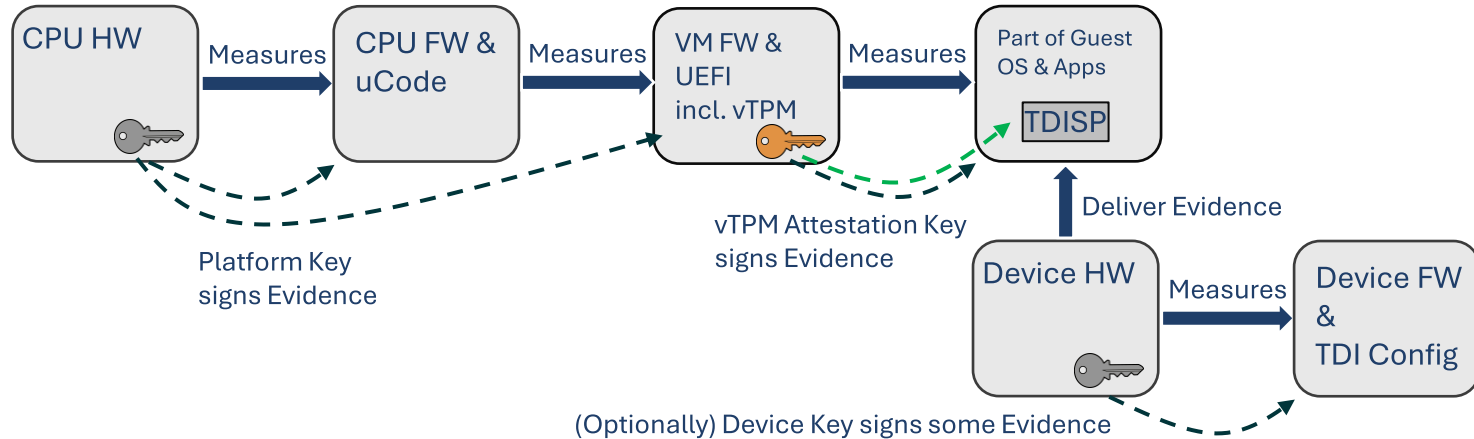
Digest  $\leftarrow$  Hash(TDISP Device ID, Device Certificate, Measurements, Interface Report)  
NVIndex\_New\_Value  $\leftarrow$  Hash(NVIndex\_Old\_Value | Digest)  
Create a new Event Log Record <NVIndex, extend event, digest>  
Append the new Event Log Record to the NVIndex Event Log file

- Every time the CVM wants to invoke a remote Verifier to perform attestation, the CVM would do:

NVIndex\_Quote  $\leftarrow$  Sign(NVIndex, vTPM Attestation Key)  
Send to remote Verifier:

- (existing) Platform Attestation Evidence and all supporting certificates
- (existing) All TDISP Device Attestation Evidence
- NVIndex\_Quote
- NVIndex Event Log file

# Binding Proposal (2/3)



- By generating the NVIndex Quote using the vTPM Attestation Key we bind the Platform Attestation chain of trust to the TDISP Attestation



# Binding Proposal (3/3)

- NVIndex needs to be defined by VM FW using Platform hierarchy
  - Platform hierarchy is under the control of the platform manufacturer (TPM) or VM FW (vTPM)
  - Platform hierarchy is needed to use TPMA\_NV\_POLICY\_DELETE attribute, which restricts everyone else (i.e. Guest) from deleting and redefining the NVIndex and its data
    - This would prevent an attacker with Guest OS kernel privileges from deceiving the remote Verifier into approving a malicious TDISP device
- Other NVIndex attributes to use:
  - TPMA\_NV\_POLICY\_READ, TPMA\_NV\_OWNERWRITE, TPM\_NT = 0x4 (extend-only), TPMA\_NV\_CLEAR\_STCLEAR (reset on reboot)
- Use Canonical Event Log Record format for Event Log format

# Open Questions

- Maintain history of all TDISP devices that were evaluated, including rejected TDISP devices and accepted TDISP devices with outdated evidence
  - Each TDISP device evidence is a few KBs in size
  - A Relying Party might use this information in its policy
    - i.e. a CVM that has previously accepted a device with vulnerable FW might have become compromised, even though the device FW got updates to a non-vulnerable version
  - Still, can cause unnecessary bloating for long lived CVMs and attestation requests to the remote Verifier
  - Need some policy over log and NVIndex reset
- Decide on location and permissions of Event Log file and NVIndex extend
  - Isolated to kernel drivers or allow applications?



Thank you!





# LINUX SECURITY SUMMIT

NORTH AMERICA