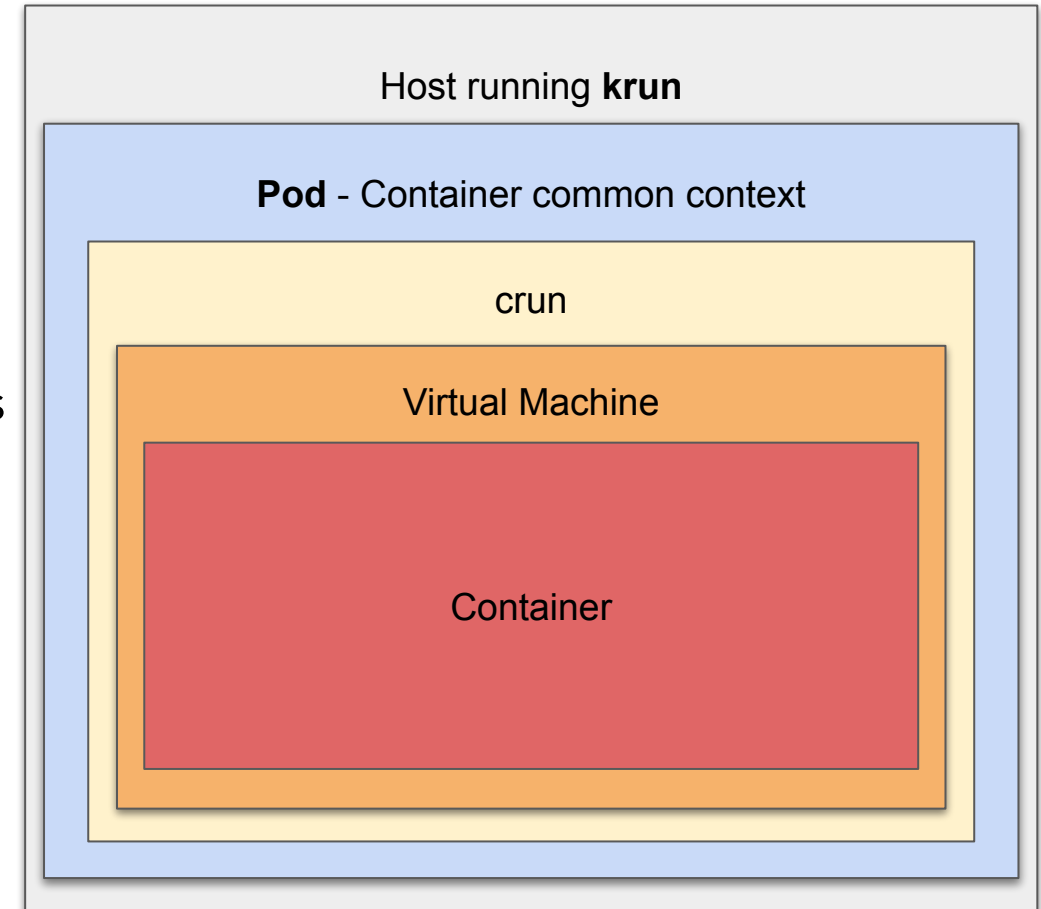


Attestable Containers with keybroker

Tyler Fanelli, Red Hat

Background: **krun** Container Runtime

- Extension of **crun** runtime with added virtualization capabilities.
- Allows containers to run inside micro-VMs, a machine type optimized for boot time and memory footprint.
 - Much smaller footprint than standard VMs, as **krun** only provides a “thin layer of virtualization” wrapping a container’s workload with a minimal amount of devices.
- Added SEV-SNP support (with attestation).



keybroker

- Simplified attestation server that re-uses the KBS protocol found in the Trustee (formerly CoCo KBS) server.
 - Experimenting with keybroker usage of the “attestation-service” (backend) of Trustee service.
- Follows RATS
- Modifications around registration and handling of registration values needed for running “confidential container images” (to be explained).

Container Images for Confidential Computing

- Users want confidence that their workloads are running confidentially.
- How can we build container images specifically with confidential computing in mind?
 - Is there a way we can “force” attestation of a container?
- Can we build container images such that attestation is required for running the container?
- Can we leverage existing container image builders to create images for confidential computing?
- How can we make the process of building “confidential container images” as simple as possible?


Component: buildah

- Tool to build OCI container images.
- Two new features:
 - **buildah build**, the main command to build OCI images, recently introduced the **--cw** (confidential workload) flag to build OCI images with **krun** VM attestation in mind
 - **buildah mkcw** command, which takes an existing container image and converts it into a “confidential workload” image.
- Both commands:
 - Build “confidential workload” images.
 - Register build information with an attestation server.



buildah Confidential Workload Images

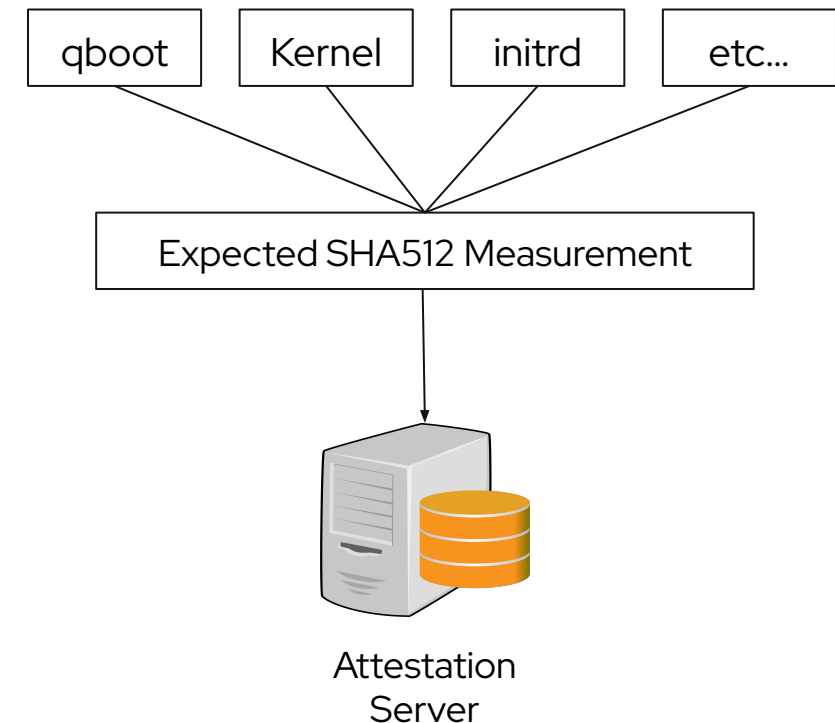
- Expected to be run with **krun** runtime.
- Contain two main disks expected by **krun**.
 - Root disk
 - Contains application code/data.
 - LUKS encrypted, with key stored in attestation server (more later).
 - Data disk
 - Contains data needed for attestation, such as the URL of the remote attestation server needing to communicate with.
- Since application (encrypted in root disk) cannot be ran until unlocked, attestation is required for all workloads.



buildah - - cw Image


buildah Confidential Workload Registration

- Buildah measures each VM component that will be used to run workload and computes a hash.
 - qboot, kernel, initrd, etc..
- Hash is registered with the attestation server as the **expected launch measurement**.
- URL of attestation server is inserted into image data disk.
- NOTE: All registration of reference values, other data needed for attestation done at build time of **container image**.
 - Any user of image must attest with attestation server that the image was **built to attest with**.
 - Whoever runs the image **must** attest it.



Traditional RATS Server Approaches


- **Every** client must attest according to a global attestation policy.
- Evidence must be presented in a form that attestation policy can interpret.
- Reference values must be submitted with policy.



Problems?

- Not flexible, all reference values/policies must be supplied in one global attestation policy.
- Doesn't fit with the buildah confidential container image model.
 - No assumptions made about the workload (reference values, TEE arch, policy) to be attested.
 - Users cannot submit their own policies when building confidential container images.
- All reference values and policies are workload-specific and aren't generated until build time.
- Would rather have a registration interface that is client-specific, where users could register their own reference values/policies and attest using them.
 - More flexible.

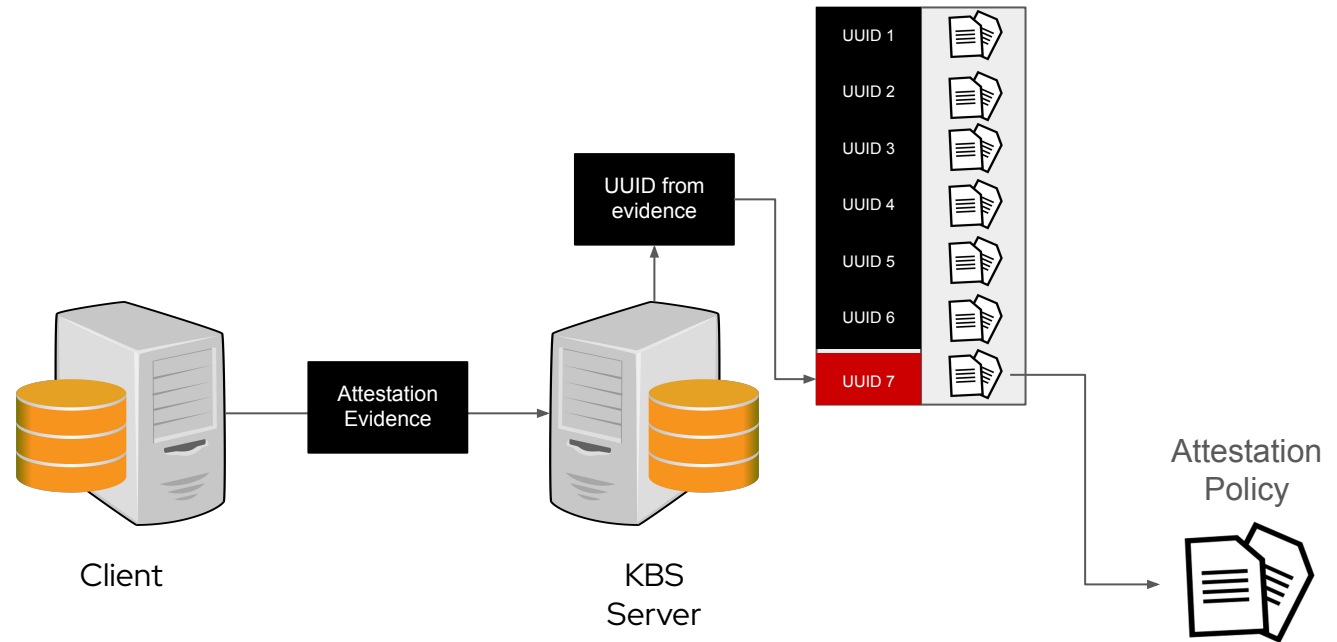
keybroker Registration



- keybroker maintains a registration endpoint, handling POST requests to register clients for later attestation.
- User registers attestation reference values and policy.
- Server generates a UUID and stores the policy in a map with the UUID as a key.
- UUID is returned to buildah. Buildah inserts ID into container image's data disk for attestation.

keybroker Attestation

- Guest attests, presenting the UUID that buildah received at registration.
- Server parses this ID from evidence.
- Server uses this ID to index into the reference value map.
- Server retrieves the policy corresponding to the client and uses reference values/policy to attest.

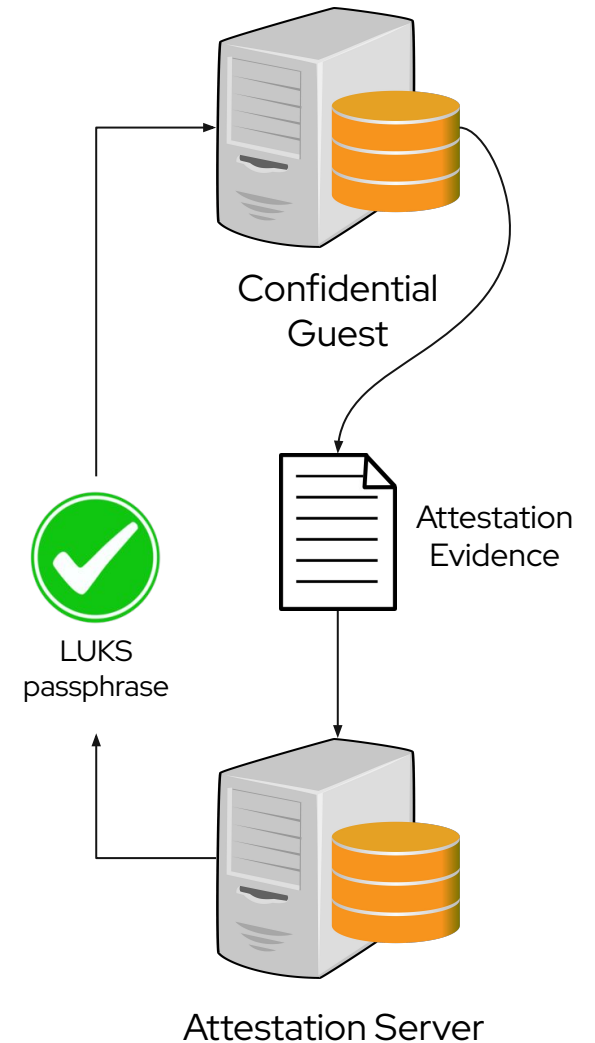


Attesting and Running Confidential Container Images

- We now have container images that must be attested to run.
- How can we use an existing container runtime to do this attestation, unlock the root disk, and run the container's contents virtualized?
- Can we hide all of this from the user, such that it's completely transparent?
- Can we hide the fact the container is running virtualized from the user?

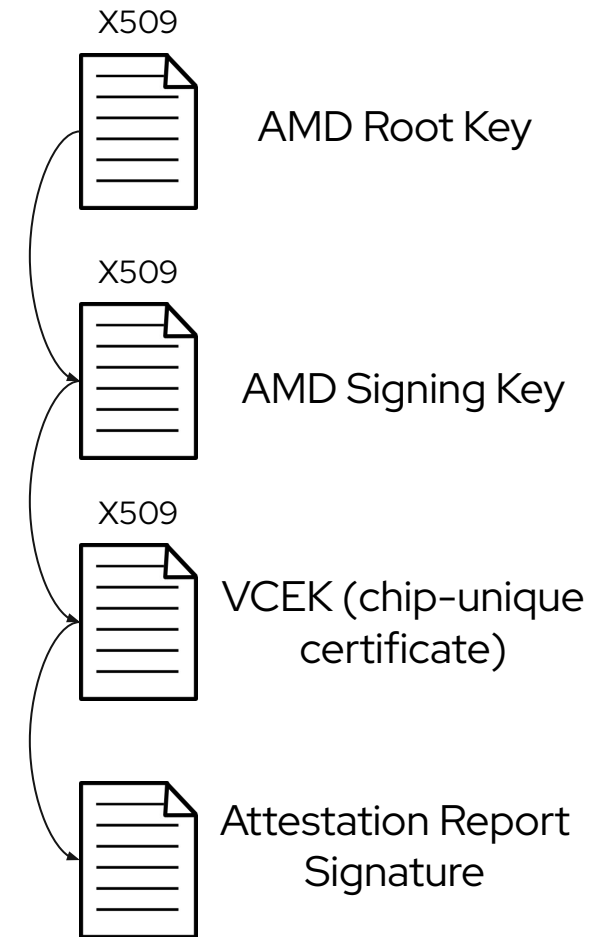
krun Attestation

- Each TEE architecture allows guests to request an attestation report from the secure processor.
 - Launch measurements, identification information, etc...
- **krun** fetches and sends an attestation report to the attestation server URL (read from data disk).
- Attestation server completes attestation and reports results.
- If successful, server sends VM the LUKS passphrase to unlock the root disk and begin running the application.



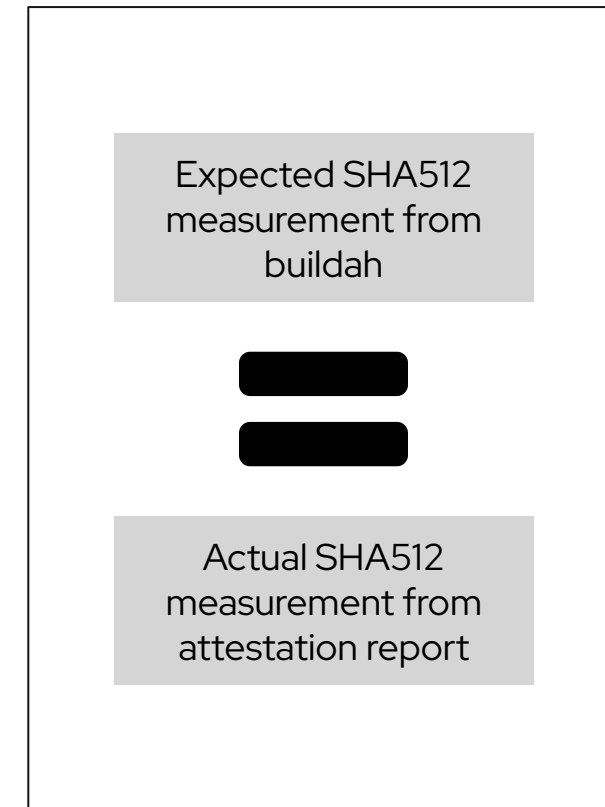
Attestation Server: Hardware Validation

- Verify that we're running on authentic TEE hardware from chip supplier.
- Each attestation report contains a signature that can be traced back to the chip supplier's root of trust.
 - SEV-SNP given as example. Certificate chain traced back to AMD root of trust.
- Cryptographically proves that the attestation report is from an authentic TEE processor.
- Host cannot lie about running confidentially.



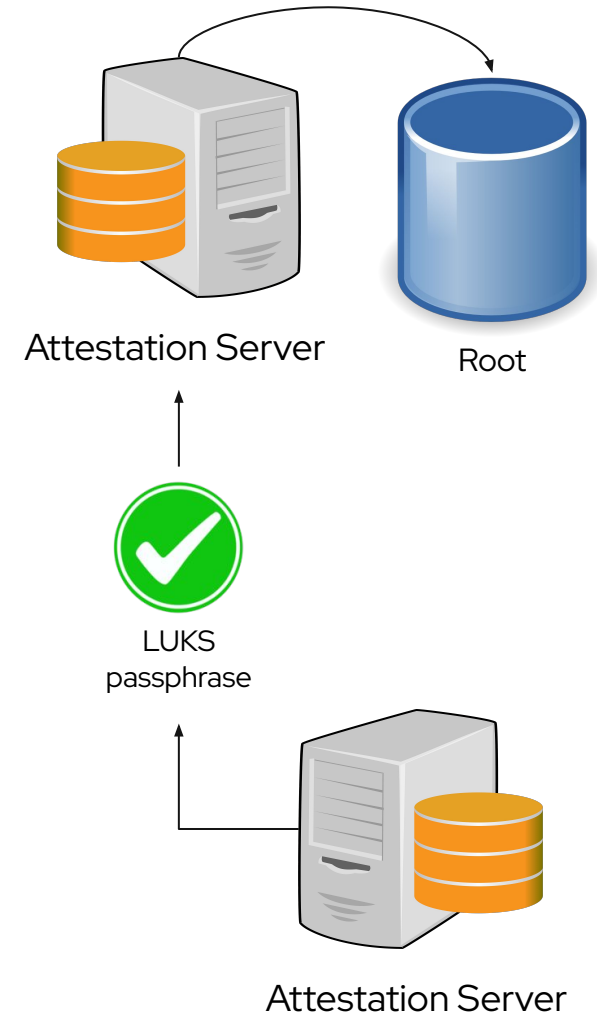
Attestation Server: Software Validation

- Recall that **buildah build --cw** previously registered an expected launch measurement with the attestation server.
- The (now authenticated) attestation report contains a hash of all components encrypted by the secure processor (i.e. the **actual launch measurement**)
- Must compare the expected launch measurement with the actual launch measurement.
- If not equal, either:
 - Not all VM pages were encrypted by secure processor.
 - Extra pages were mapped into VM guest memory.
- Either option may jeopardize the confidentiality of the workload, so attestation fails if expected and actual launch measurements do not match.



krun VM: Successful Attestation

- On successful attestation, attestation server sends LUKS passphrase to VM.
- VM uses LUKS passphrase to unlock root disk and begin running application.
- Cryptographically proven that our workload (and only our workload) is running confidentially on the host system.

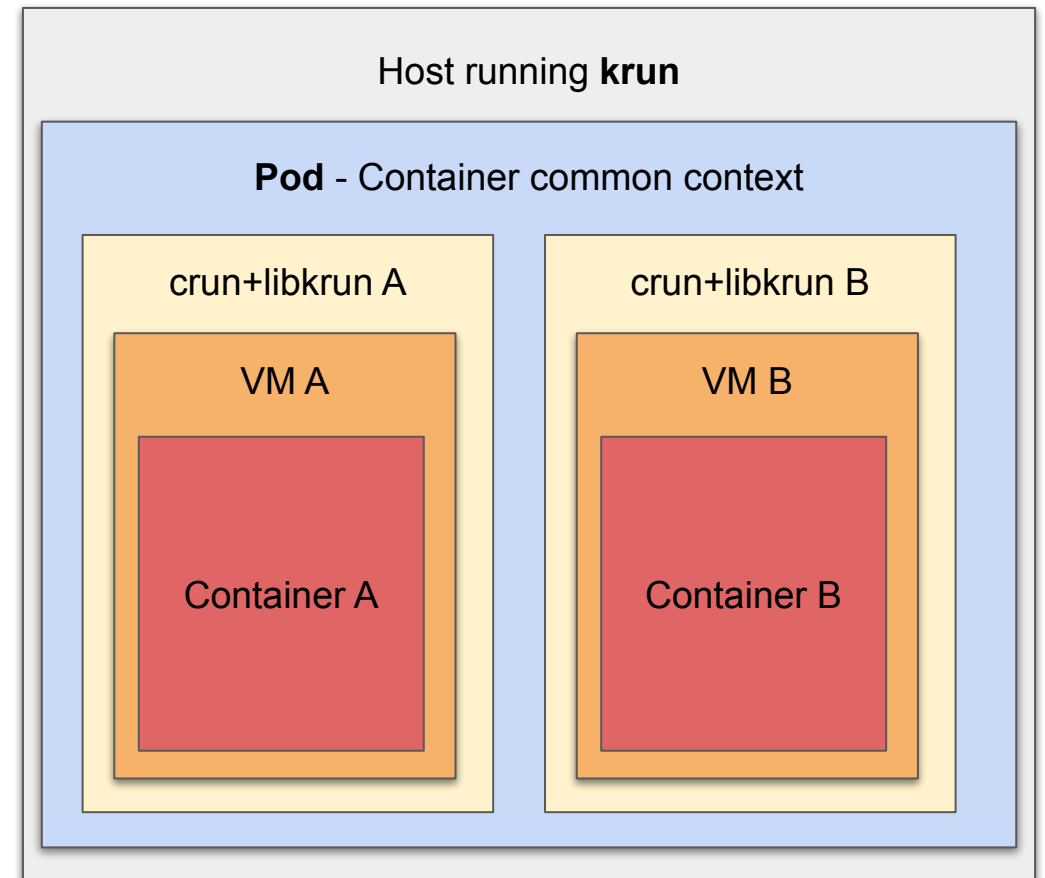
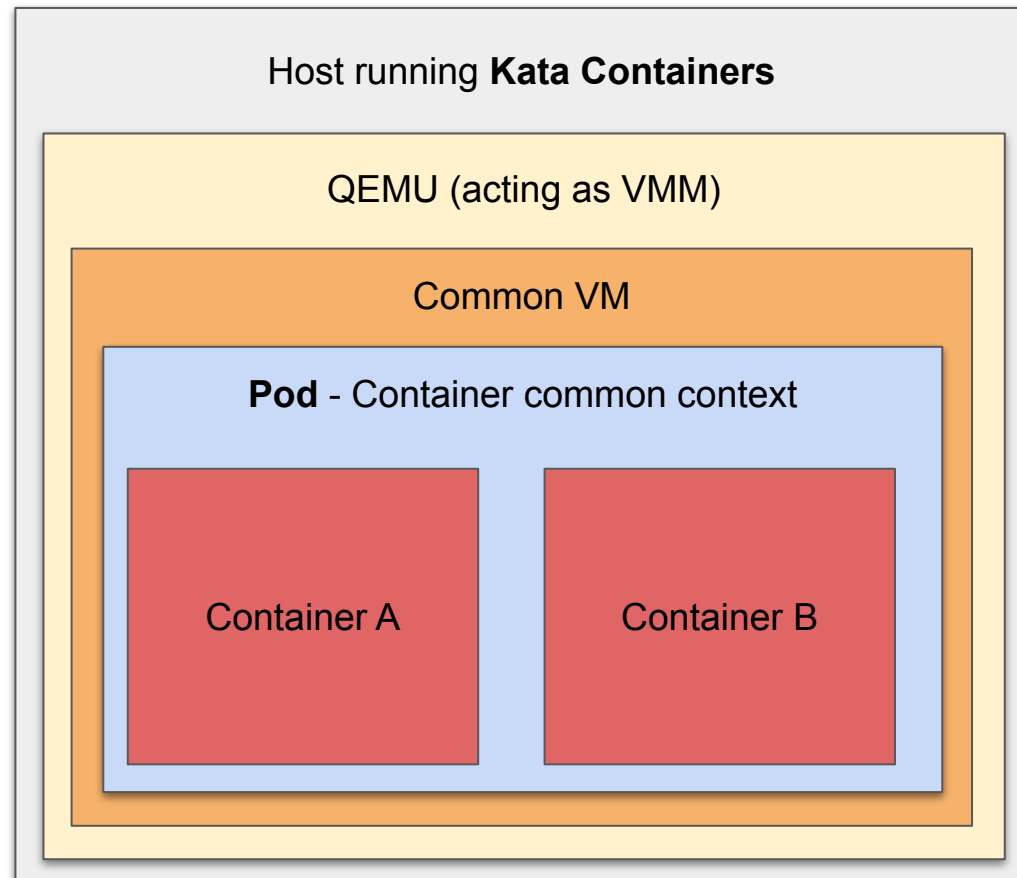


SEV-SNP Demo

<https://asciinema.org/a/653719>

Questions?

Aside: krun Runtime vs. Kata Containers



krun Runtime vs. Kata Containers

- Kata Containers advantages
 - Supports most container workloads.
 - Built on mature components (QEMU).
 - Private environment between containers in the same Pod.
- krun advantages
 - Simpler, fewer moving pieces.
 - Doesn't need to grow/shrink the VM.
 - Smaller attack surface.
 - Lower per-VM memory footprint.

Thanks!